THE UNIVERSITY OF WESTERN AUSTRALIA

Predicting Housing Prices in Perth

Presented by

Bharat Guturi

Project Overview and Source of Data

For someone who is contemplating to buy a house, this Project aims to predict the house prices in Perth based on the available historical data available.

The Data is sourced from:

https://www.kaggle.com/datasets/syuzai/perth-house-prices

Architecture diagram

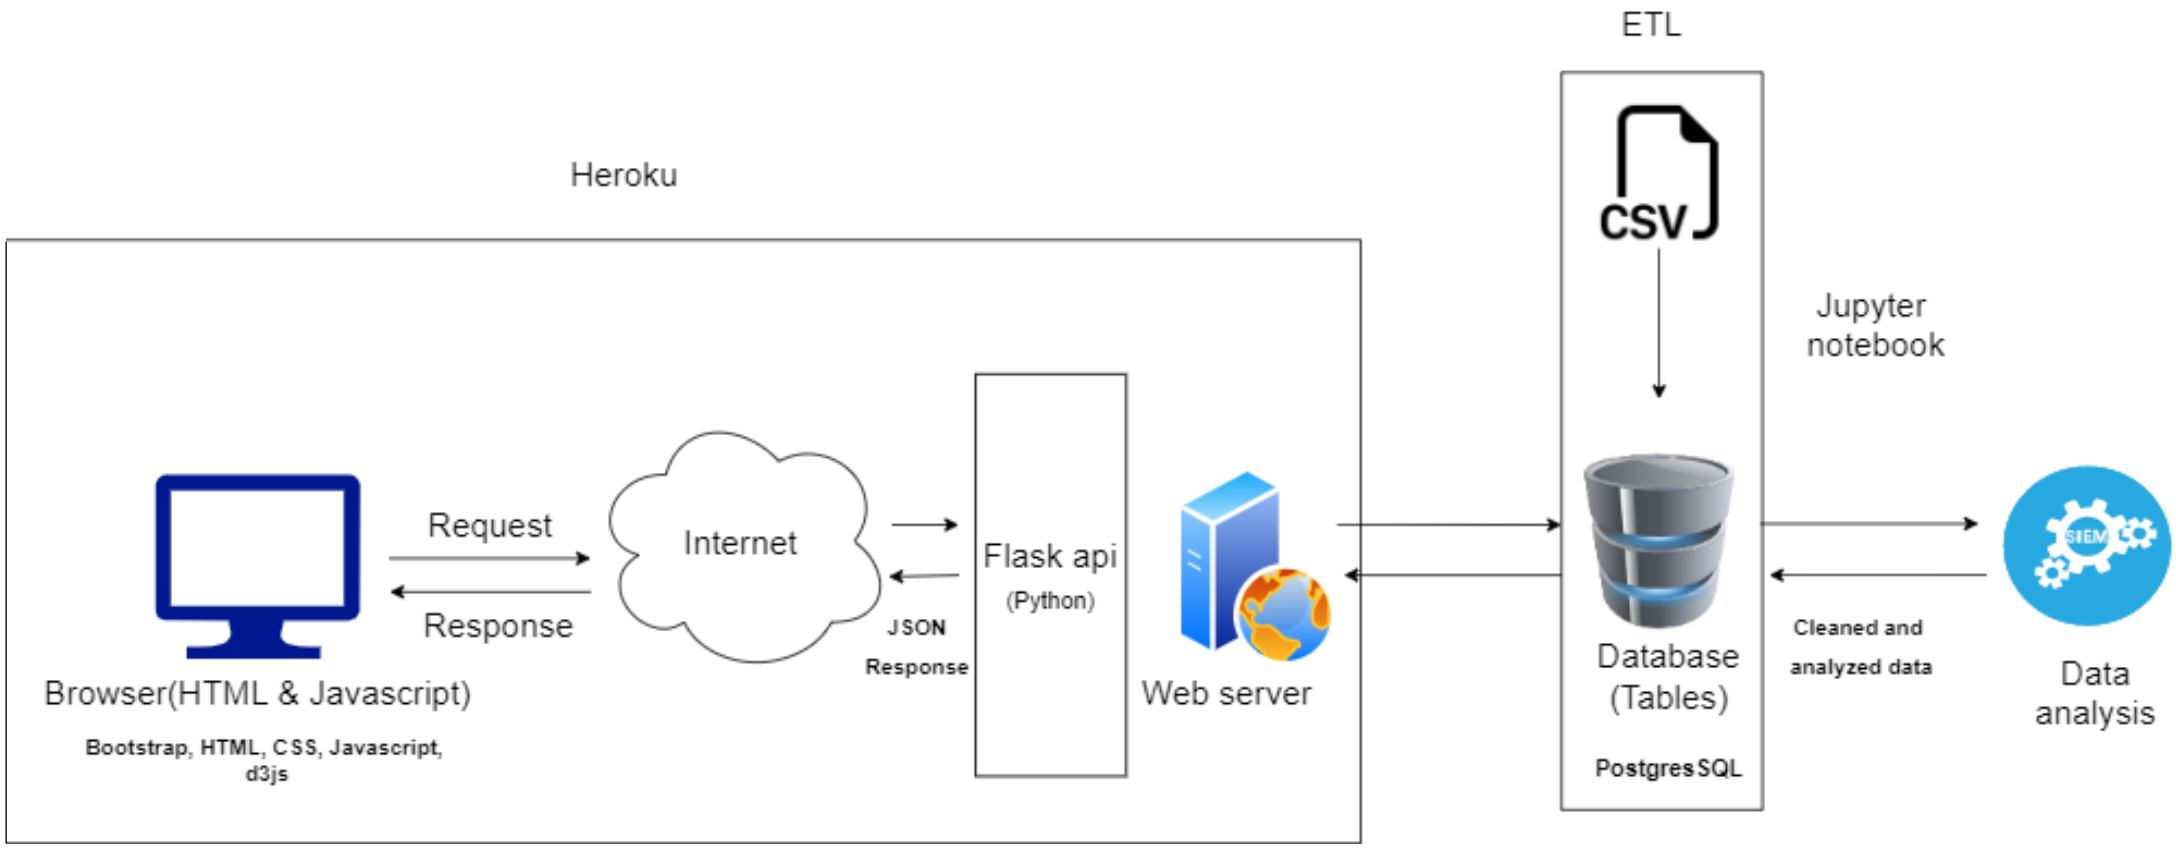

THE UNIVERSITY OF WESTERN AUSTRALIA

Extract, Transform & load

1. Importing Dependencies
2. Storing CSV into Dataframe

	ADDRESS	SUBURB	PRICE	BEDROOMS	BATHROOMS	GARAGE	LAND_AREA	FLOOR_AREA	BUILD_YEAR	CBD_DIST	NEAREST_STN	NEAREST_STN_D
0	1 Acorn Place	South Lake	565000	4	2	2.0	600	160	2003.0	18300	Cockburn Central Station	1
1	1 Addis Way	Wandi	365000	3	2	2.0	351	139	2013.0	26900	Kwinana Station	4
2	1 Ainsley Court	Camillo	287000	3	1	1.0	719	86	1979.0	22600	Challis Station	1
3	1 Albert Street	Bellevue	255000	2	1	2.0	651	59	1953.0	17900	Midland Station	3
4	1 Aman Place	Lockridge	325000	4	1	2.0	466	131	1998.0	11200	Bassendean Station	2

3. Connecting to local database

```
load_dotenv()
protocol = 'postgresql'
username = os.environ.get('db_Username')
password = os.environ.get('db_Password')
host = 'localhost'
port = 5432
database_name = 'housing_db'
rds_connection_string = f'{protocol}://{username}:{password}@{host}:{port}/{database_name}'
engine = create_engine(rds_connection_string)
insp = inspect(engine)
```

4. Loading csv data using dataframe to SQL

```
prices_df.to_sql(name='Perth_housing',if_exists = 'append', con=engine, index=False)
```

Data Analysis

Data Cleaning:

Data Columns and datatypes

```
PRICE                          int64
BEDROOMS                       int64
BATHROOMS                      int64
GARAGE                         int64
LAND_AREA                      int64
FLOOR_AREA                     int64
BUILD_YEAR                     int64
CBD_DIST                       int64
NEAREST_STN_DIST               int64
DATE_SOLD             datetime64[ns]
POSTCODE                       int64
LATITUDE                     float64
LONGITUDE                    float64
NEAREST_SCH_DIST             float64
NEAREST_SCH_RANK               int64
dtype: object
```

Steps:

1. Removal of rows containing null values in 'Nearest_SCH_RANK'
2. Changing the 'DATE_SOLD' column to 'MONTH_SOLD' and 'YEAR_SOLD'.

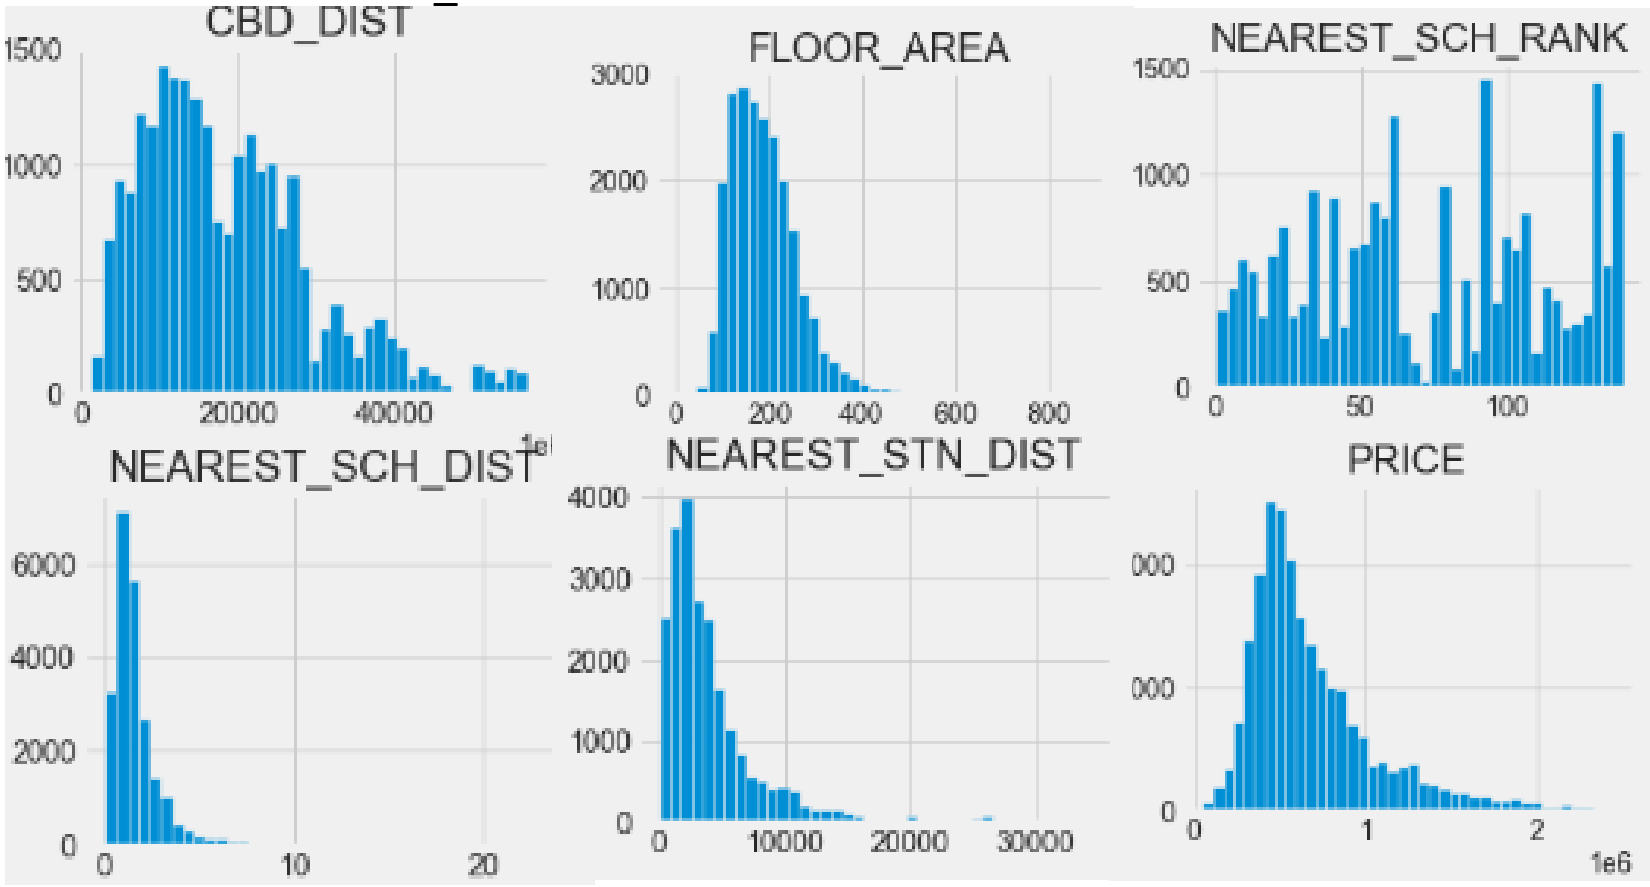

THE UNIVERSITY OF WESTERN AUSTRALIA

Box Plots

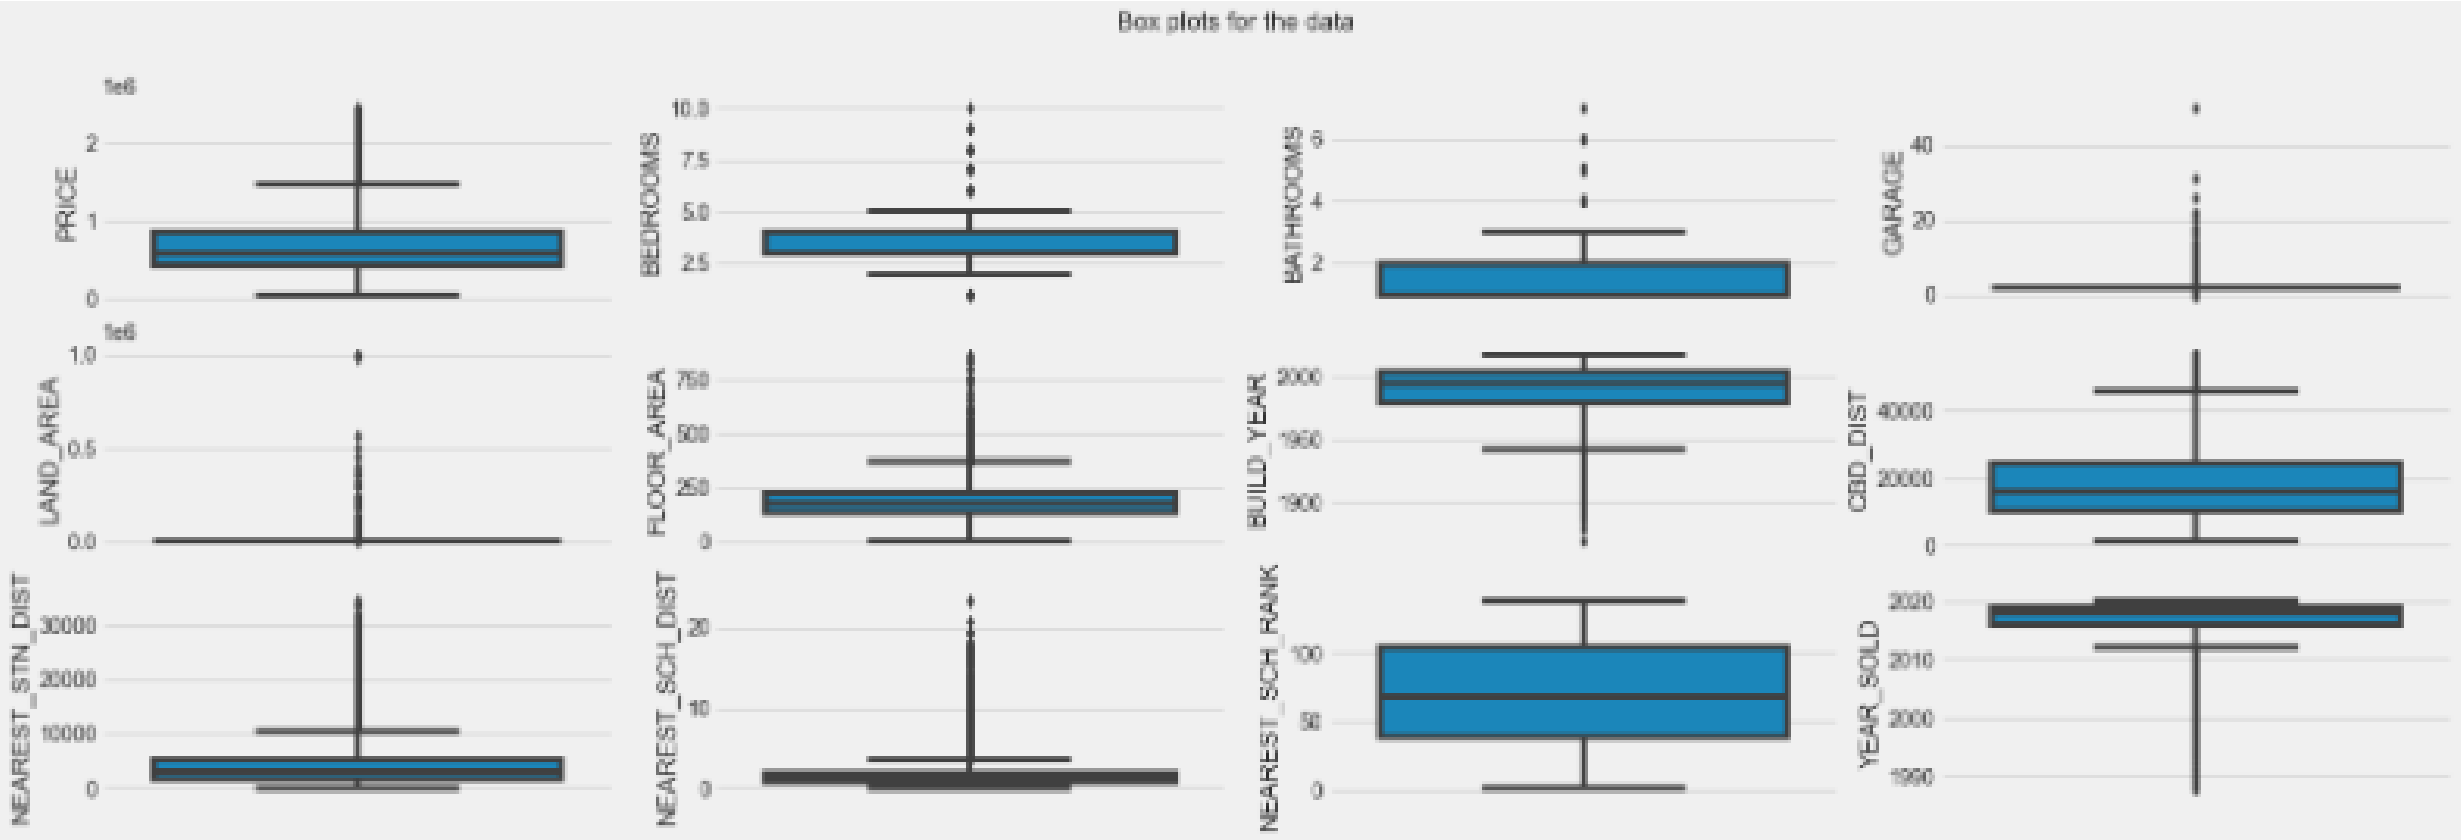

THE UNIVERSITY OF WESTERN AUSTRALIA

Filtered Data

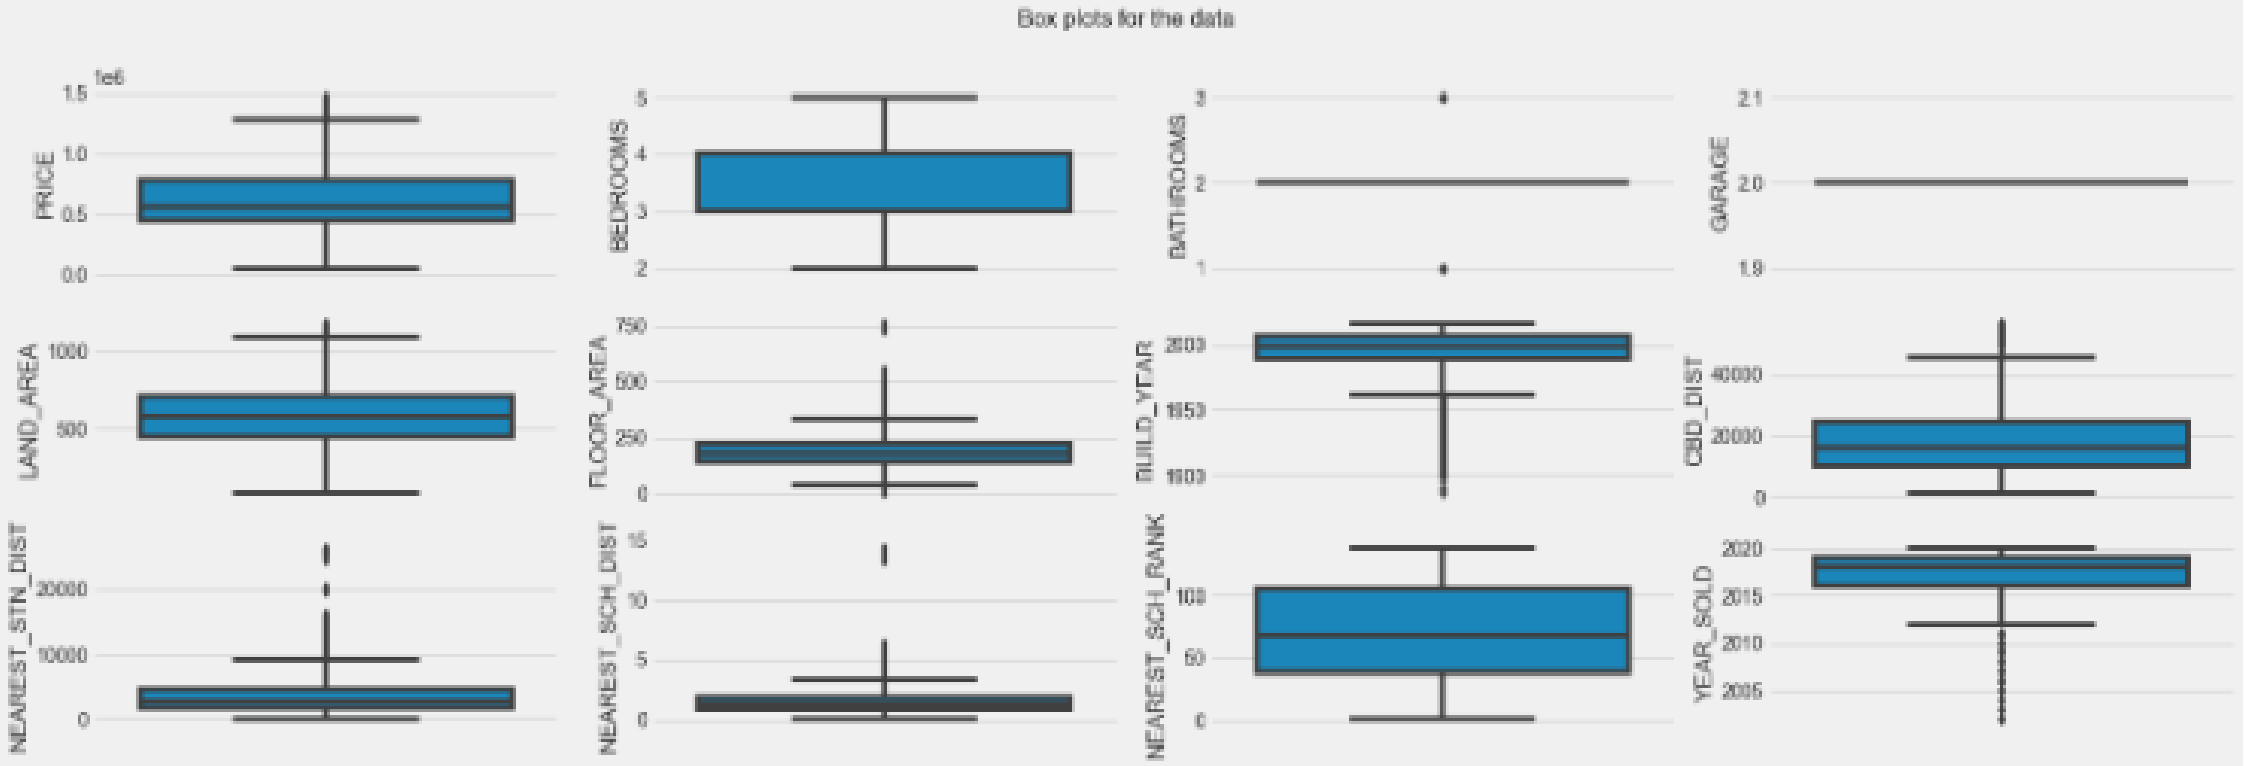

THE UNIVERSITY OF WESTERN AUSTRALIA

Correlation Matrix

```
PRICE                1.000
FLOOR_AREA           0.544
BATHROOMS            0.382
BEDROOMS             0.246
GARAGE               0.136
YEAR_SOLD            0.053
LATITUDE             0.049
LAND_AREA            0.027
MONTH_SOLD          -0.005
NEAREST_SCH_DIST    -0.057
NEAREST_STN_DIST    -0.142
BUILD_YEAR          -0.176
POSTCODE            -0.206
LONGITUDE           -0.227
CBD_DIST            -0.393
NEAREST_SCH_RANK    -0.462
Name: PRICE, dtype: float64
```

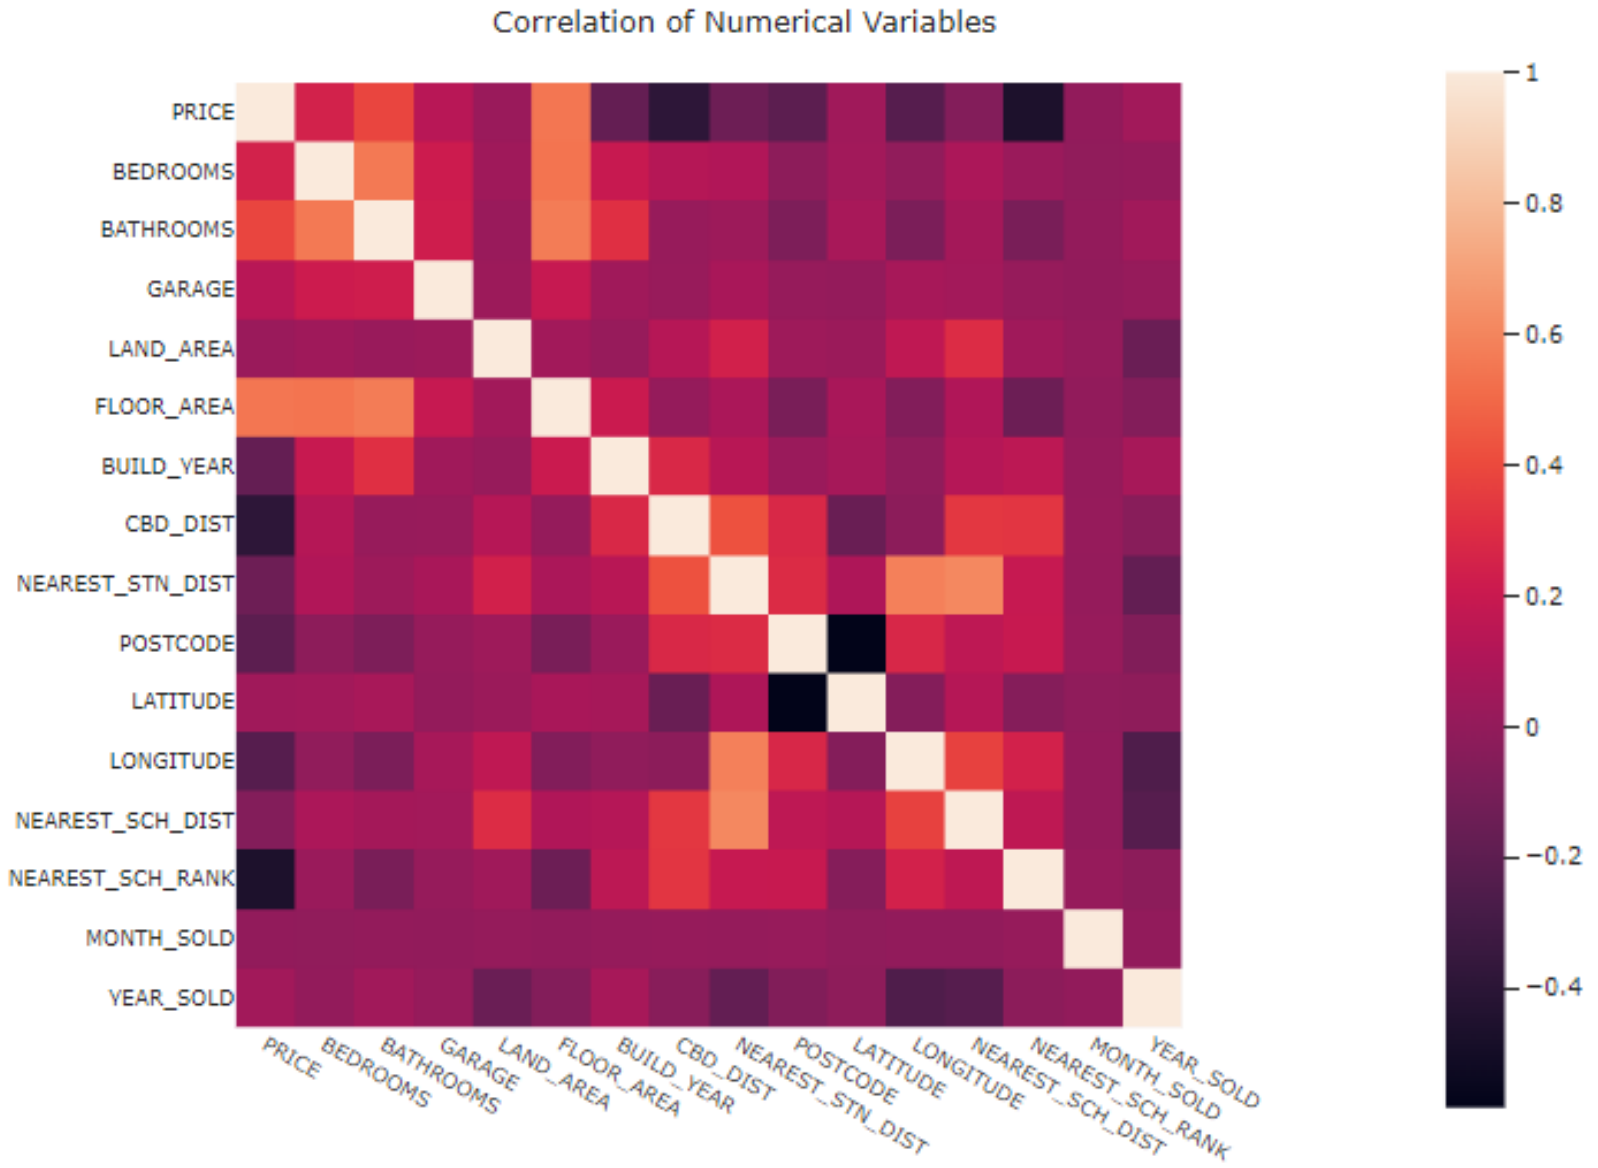

THE UNIVERSITY OF WESTERN AUSTRALIA

Visualisation of Filtered Data as per the Price

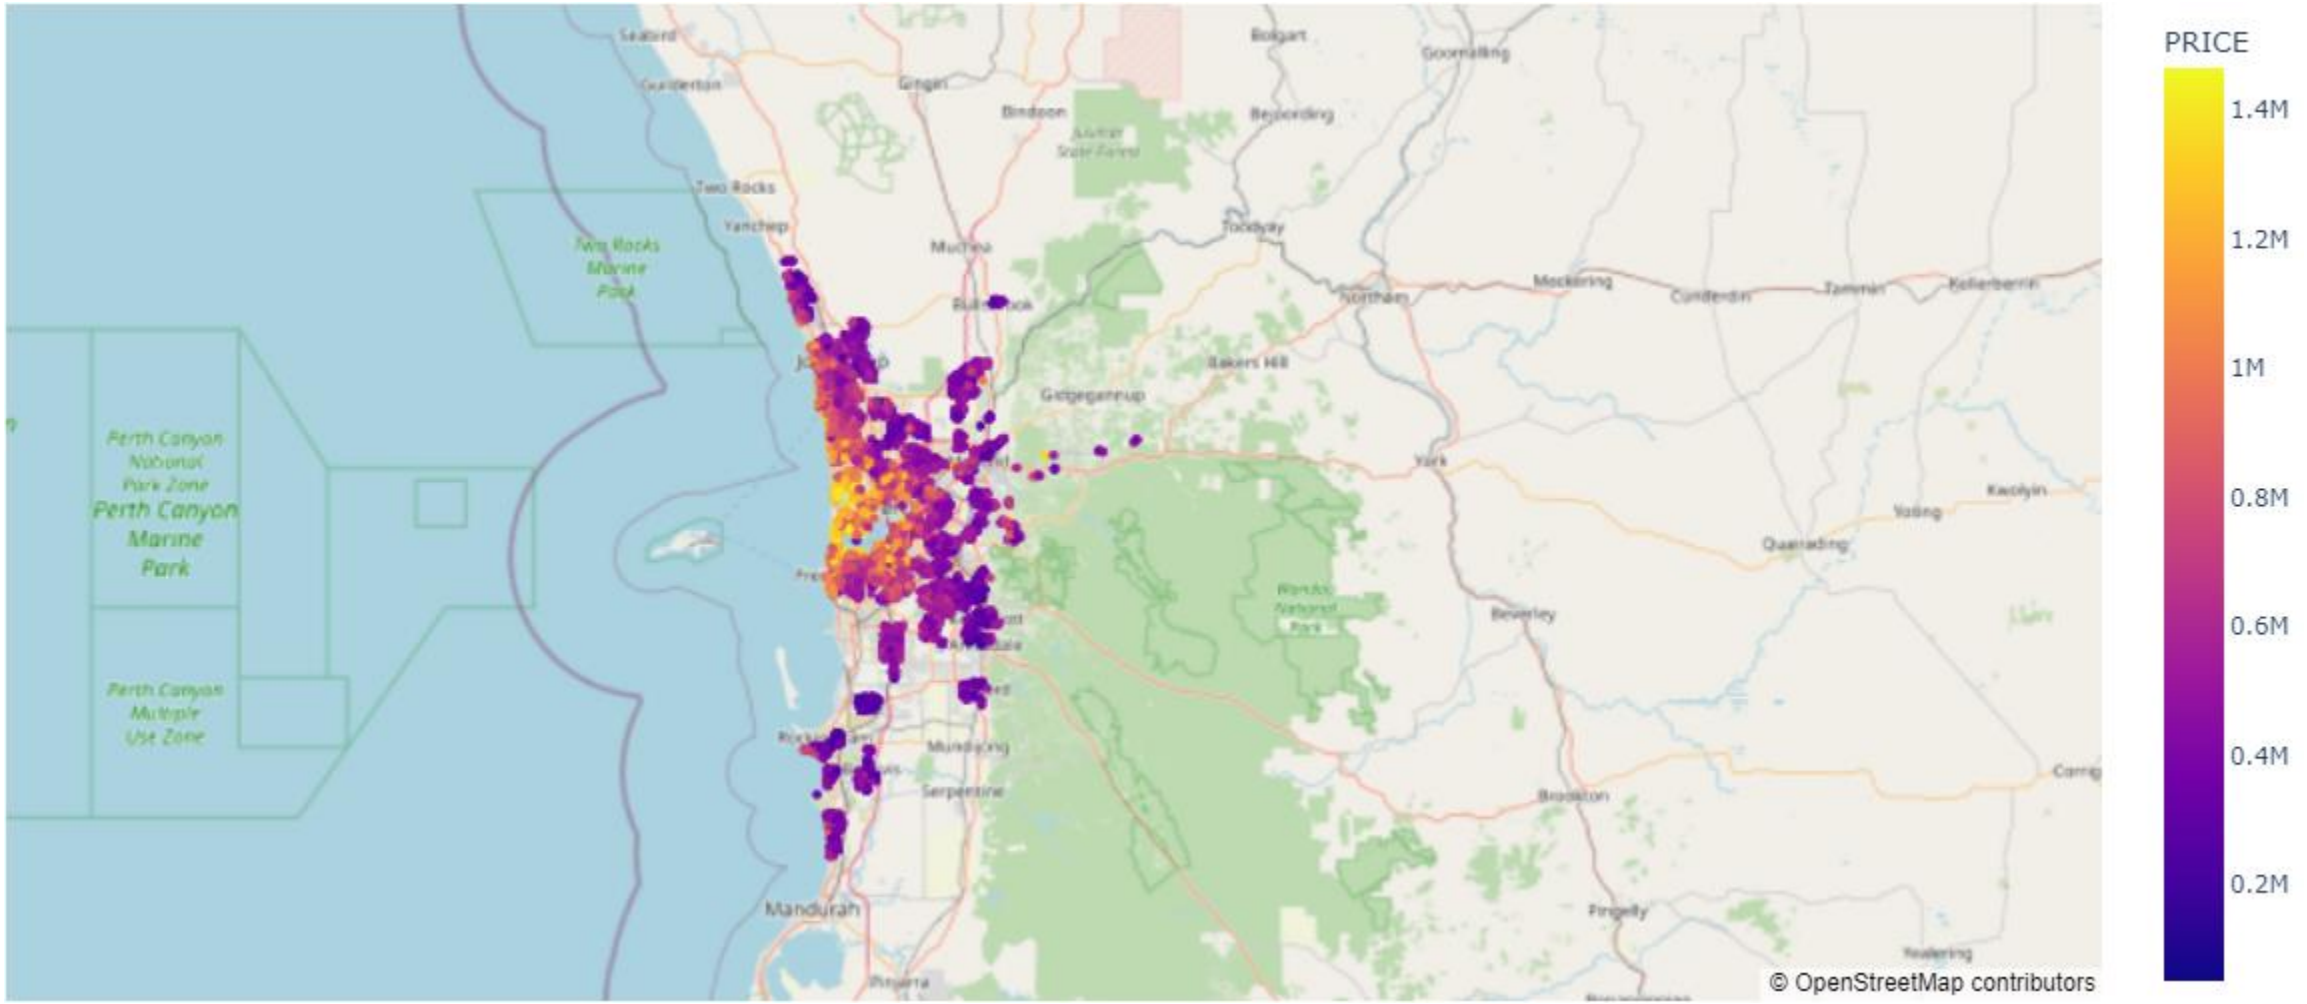

Choosing the right Algorithm

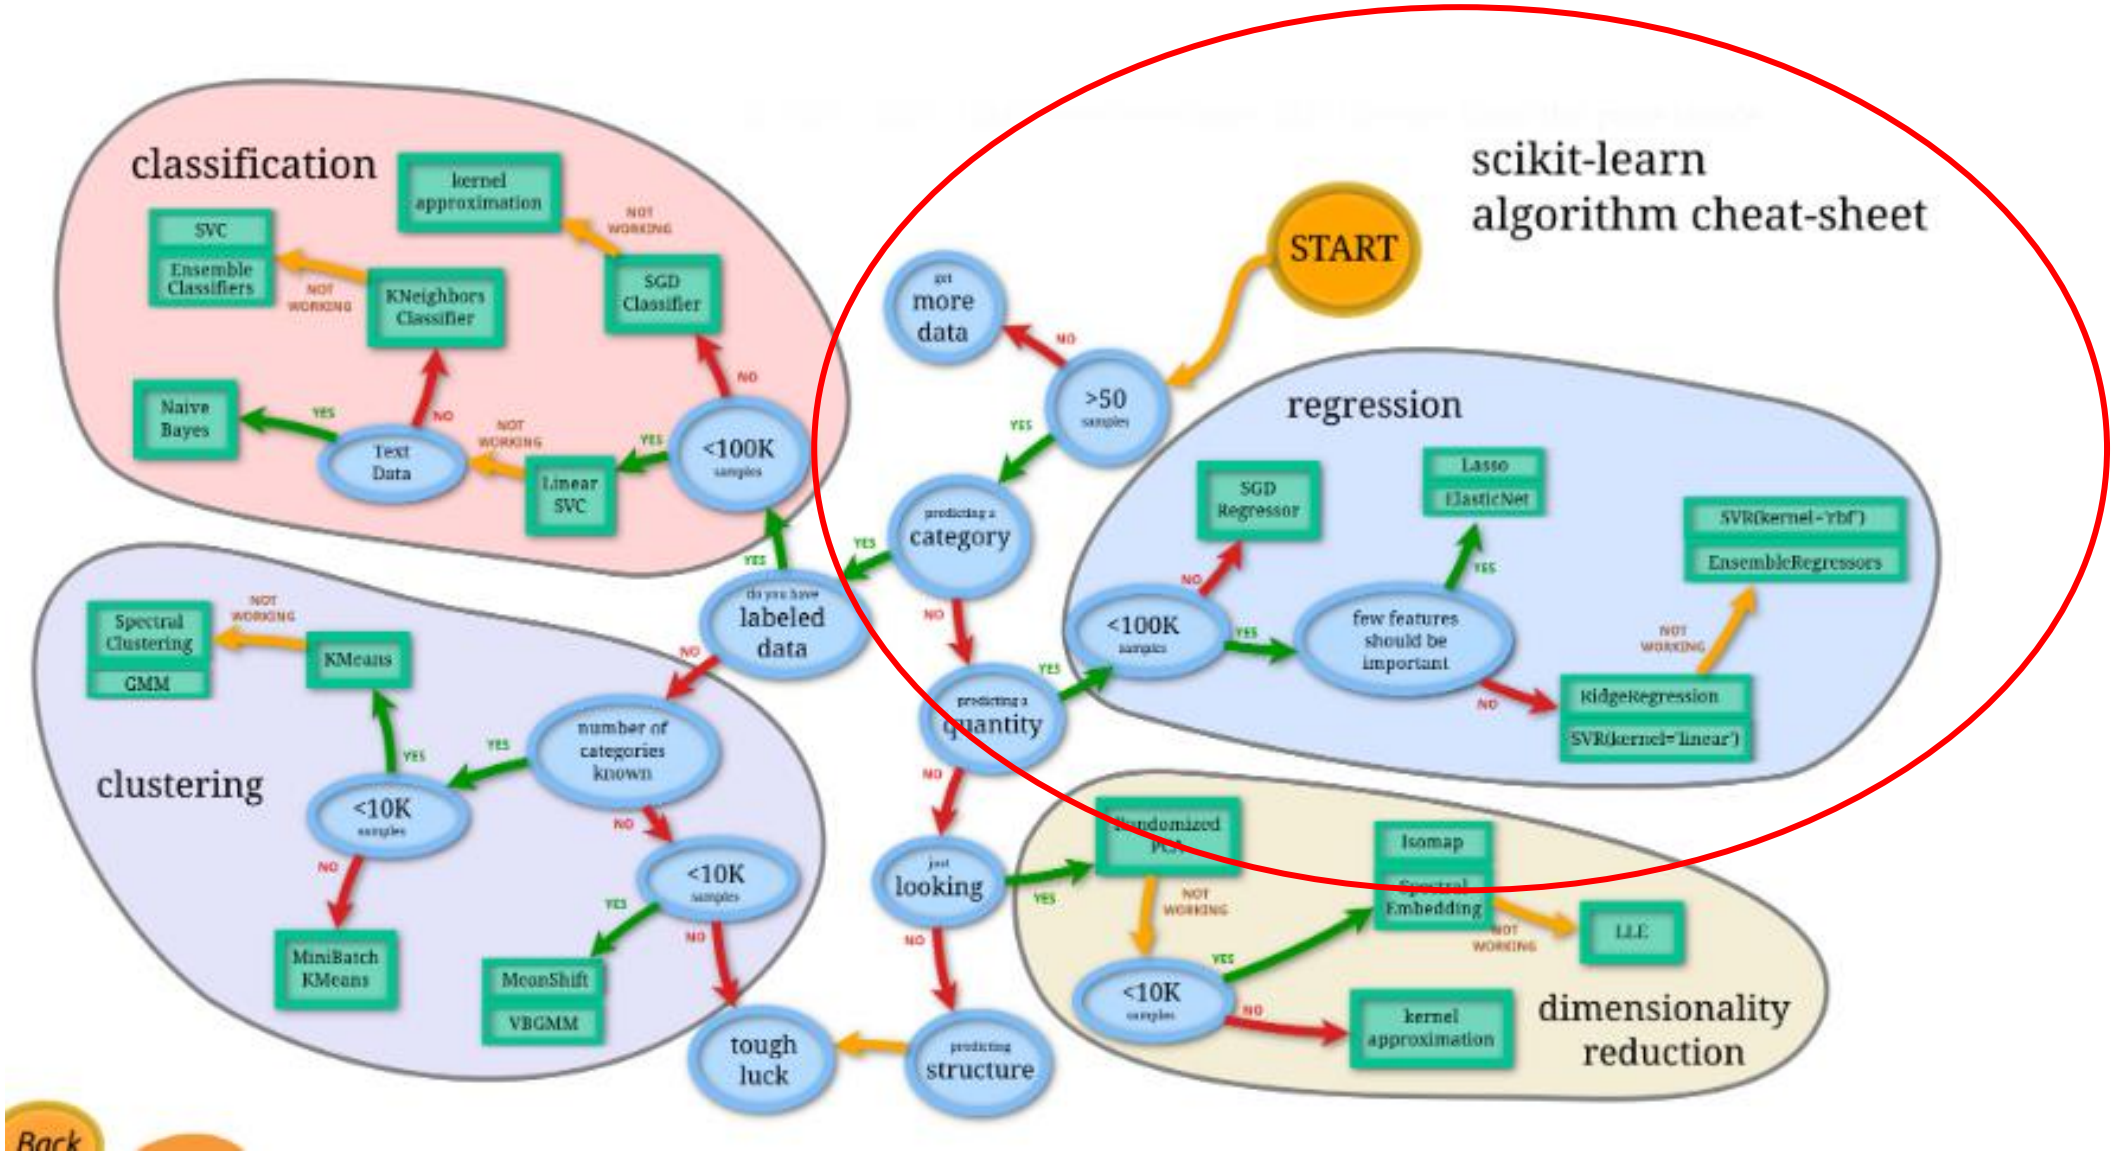

THE UNIVERSITY OF WESTERN AUSTRALIA

Resulting Accuracy of various Models

	Model	R2 Score	Accuracy
0	RandomForestRegressor	0.851	87.652
5	DecisionTreeRegressor	0.711	83.590
3	Ridge	0.701	81.247
1	BayesianRidge	0.701	81.253
2	ElasticNet	0.121	63.371
4	SVR	-0.078	66.671

Random Forest Regressor

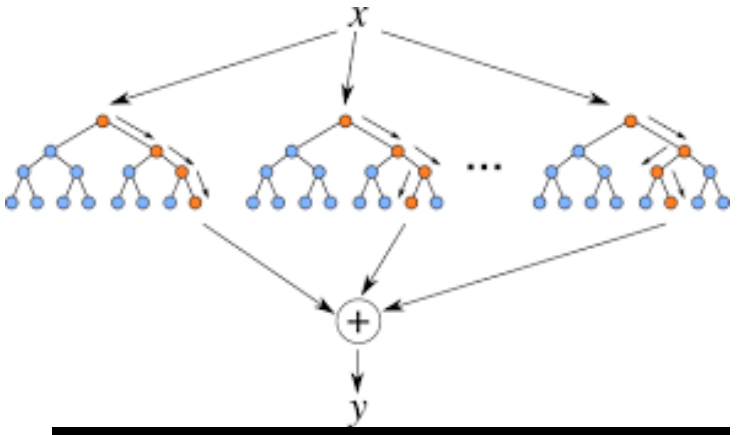

Ridge Regressor

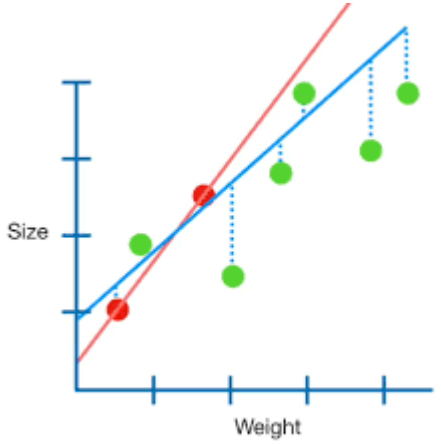

Decision Tree Regressor

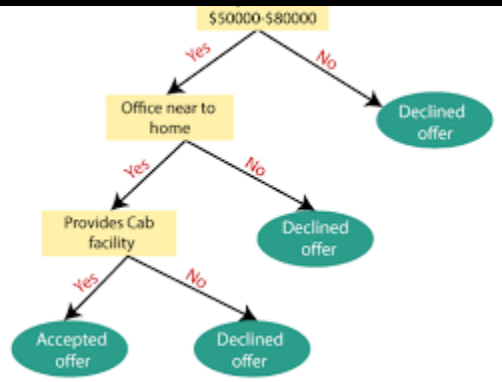

Bayesian Ridge Regressor

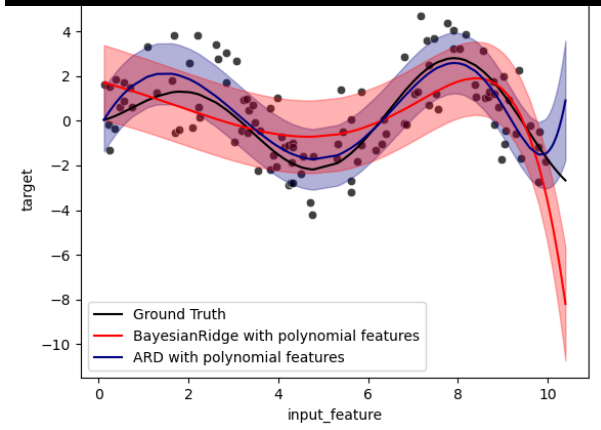

SVR

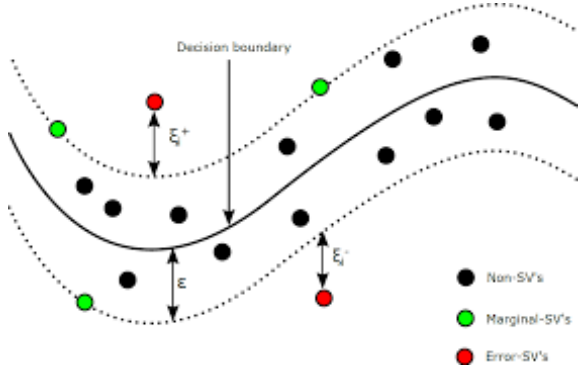

THE UNIVERSITY OF WESTERN AUSTRALIA

Hyperparameter Tuning

```
# Use the random grid to search for best hyperparameters
# First create the base model to tune

# Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, num = 10)]
# Number of features to consider at every split
max_features = ['auto', 'sqrt']
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
max_depth.append(None)
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4]
# Method of selecting samples for training each tree
bootstrap = [True, False]

# Create the random grid
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'bootstrap': bootstrap}

rf = RandomForestRegressor(random_state = 42)
# Random search of parameters, using 3 fold cross validation,
# search across 100 different combinations, and use all available cores
rf_random = RandomizedSearchCV(estimator=rf, param_distributions=random_grid,
                               n_iter = 10, scoring='neg_mean_absolute_error',
                               cv = 3, verbose=2, random_state=42, n_jobs=-1,
                               return_train_score=True)

# Fit the random search model
rf_random.fit(X_train_scaled, y_train);
```

Best Parameters

```
rf_random.best_params_
```

```
{'n_estimators': 400,
 'min_samples_split': 10,
 'min_samples_leaf': 1,
 'max_features': 'sqrt',
 'max_depth': 60,
 'bootstrap': False}
```

Output

	Model	R2 Score	Accuracy
0	best_random_forest	0.711	83.590

THE UNIVERSITY OF WESTERN AUSTRALIA

Web Application

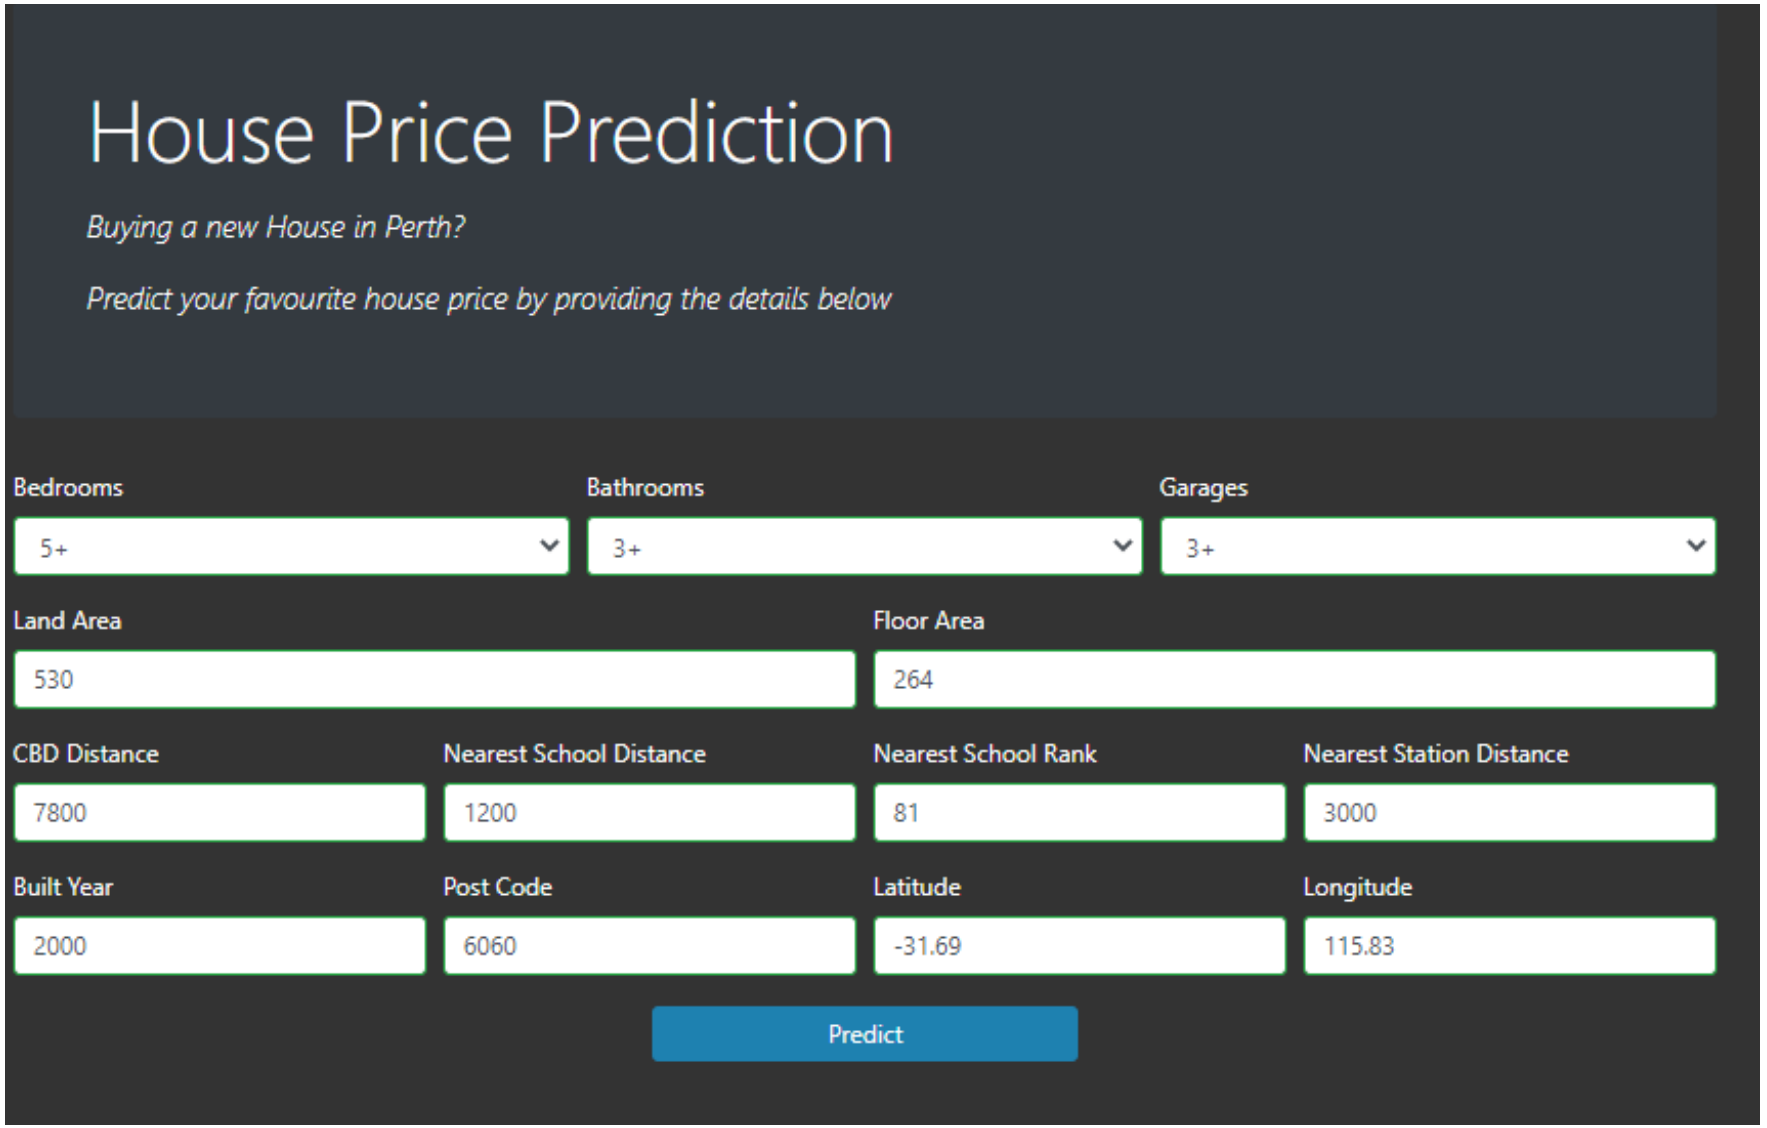

Testing

Case Study: https://www.realestate.com.au/property-house-wa-tuart+hill-140995332

Predicted House Price

913930

Close

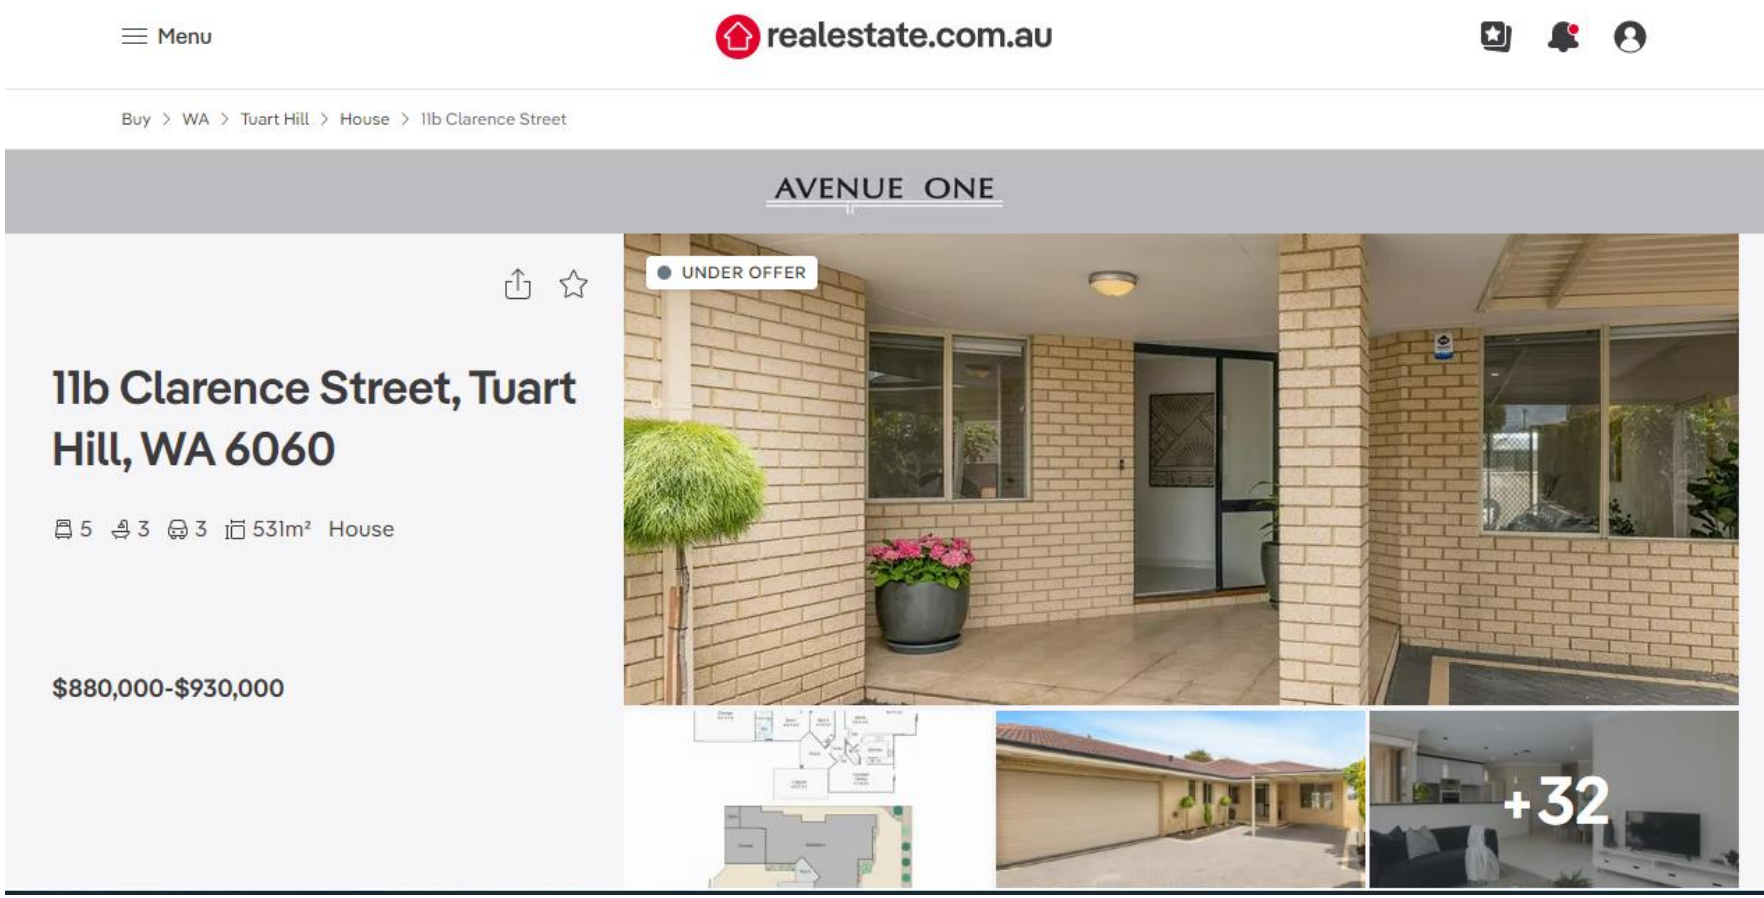

THE UNIVERSITY OF WESTERN AUSTRALIA

Thank you